



A. Kaestner :: Paul Scherrer Institut

Image enhancement

Preparing the images for analysis

Outline

- 1 Introduction
- 2 Noise and Artifacts
- 3 Basic filtering
- 4 Scale spaces
- 5 Verification
- 6 Summary

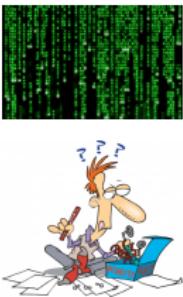
Introduction

Introduction

3D and 4D imaging produce large amounts of data



Gigabytes...
... or even
terabytes of data

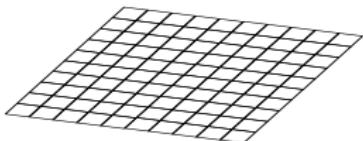


- 3D visualization
- Sample characterization
- Process parameterization
- etc

Different types of images

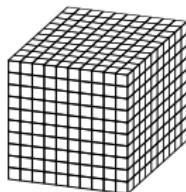
2D

- Pictures
- Radiographs
- CT slices



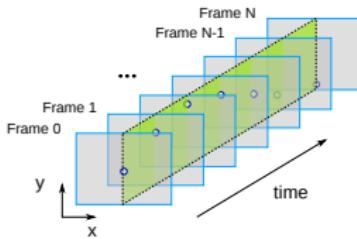
3D

- Volumes
- x, y, z
- Movies
- x, y, t



4D

- Volume movie
- x, y, z, t



Which information do you want to gain

Quantitative

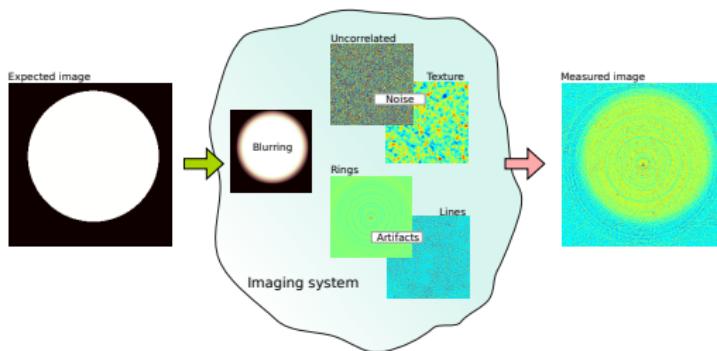
- Material composition
- Material transport

Structure

- Identify items
- Item geometry

This will affect the choice of processing methods.

Measurements are rarely perfect

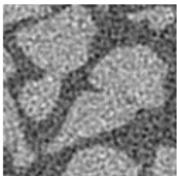


Factors affecting the image quality

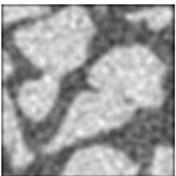
- Resolution (Imaging system transfer functions)
- Noise
- Contrast
- Inhomogeneous contrast
- Small relevant features
- Artifacts

A typical processing chain

Image
Acquisition



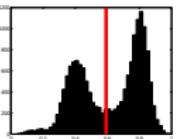
Enhancement



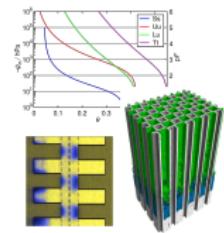
Segmentation



Post processing



Evaluation



Todays lecture will focus on the enhancement

Noise and artifacts

The unwanted information

Noise types

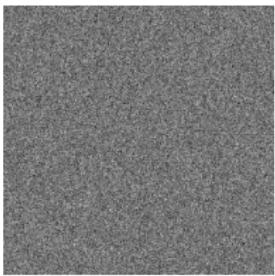
Spatially uncorrelated noise

Event noise

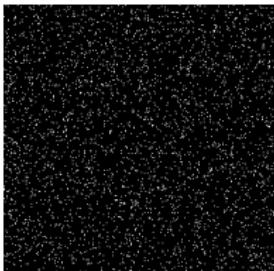
Structured noise

Noise examples

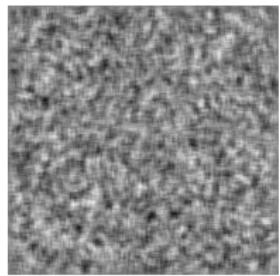
Gaussian



Salt 'n pepper



Structured



Gaussian noise

- Additive
- Easy to model
- Other distributions obtain Gaussian shape at large numbers

$$n(x) = \frac{1}{\sqrt{2\pi}\sigma} e^{-(\frac{x-\mu}{2\sigma})^2}$$

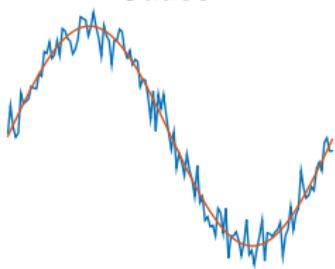
Poisson noise

- Multiplicative
- Physically correct for event counting

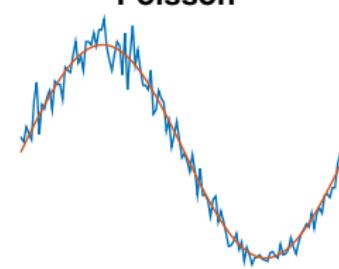
$$p(x) = \frac{\lambda^k}{k!} e^{-\lambda} x$$

Noise examples

Gauss



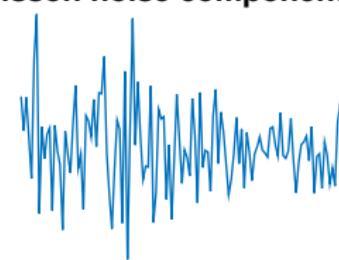
Poisson



Gaussian noise component



Poisson noise component



- A type of outlier noise
- Noise strength give as the probability of an outlier
- Additive, multiplicative, independent replacement

Example

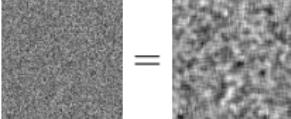
$$sp(x) = \begin{cases} -1 & x \leq \lambda_1 & x \in \mathcal{U}(0, 1) \\ 0 & \lambda_1 < x \leq \lambda_2 & \lambda_1 < \lambda_2 \\ 1 & \lambda_2 < x & \lambda_1 + \lambda_2 = \text{noise fraction} \end{cases}$$

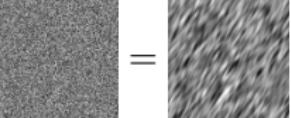
- Spatially correlated
- Example: Detector structure

Example random field models

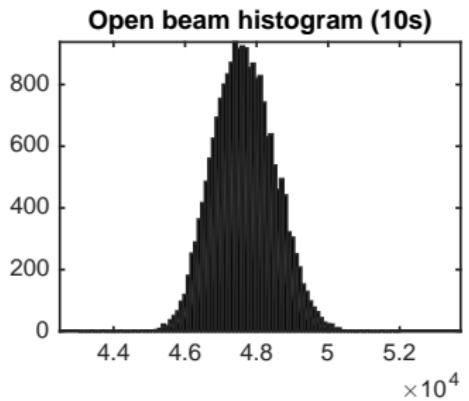
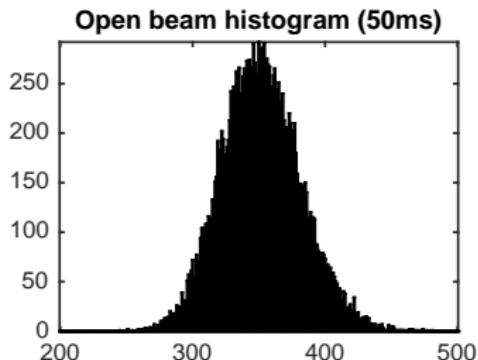
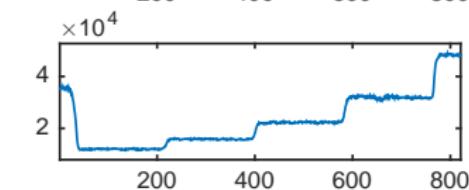
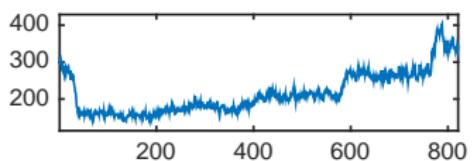
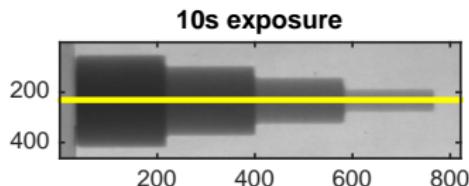
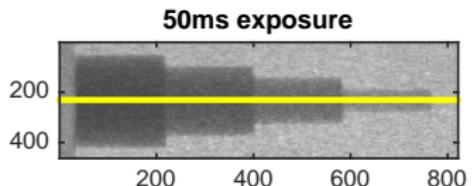
$$n(x, y) \in \mathcal{N}(\mu, \sigma)$$

$$ns = K * n \quad K = \text{convolution kernel}$$

$$U_{5 \times 5} * \quad = \quad$$


$$\quad * \quad = \quad$$


Noisy profiles



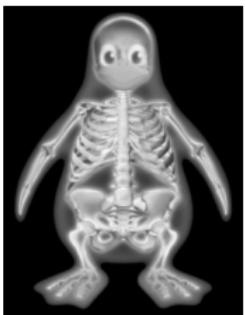
Signal to noise ratio

A metric to describe noise strength

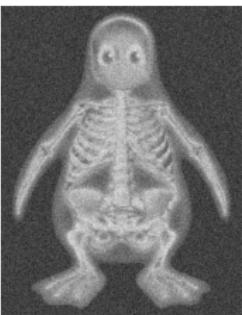
$$SNR = \frac{\mu_{image}}{\sigma_{image}} \quad (1)$$

$$SNR_{db} = 20 \log \frac{\mu_{image}}{\sigma_{image}} \quad (2)$$

- Select a region
- Compute average intensity
- Compute std deviation
- Apply eqns 1 or 2



SNR = ∞



SNR = 5



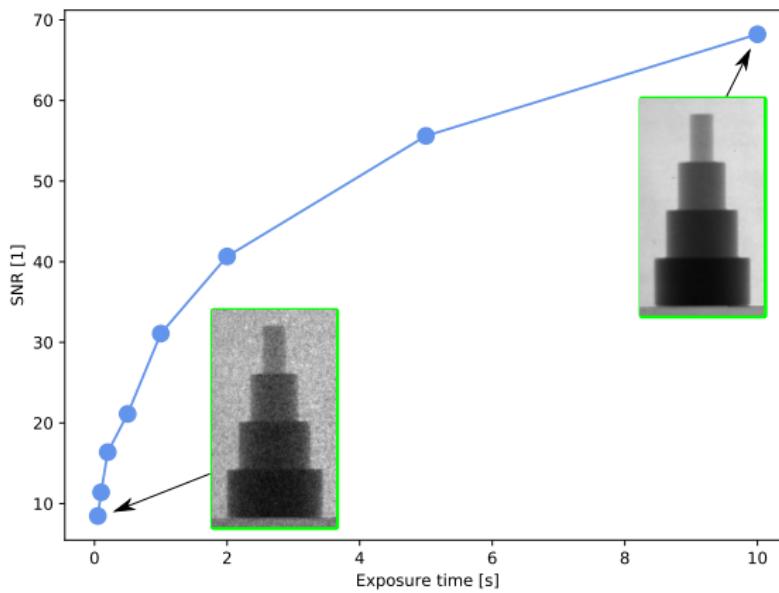
SNR = 2



SNR = 1

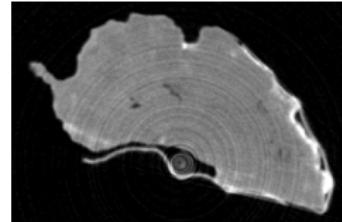
Poisson noise

- $E[Poi(\lambda)] = \lambda$ and $var[Poi(\lambda)] = \lambda \rightarrow \text{SNR} = \sqrt{\lambda}$
- Exposure time increase the number of events



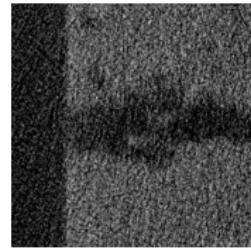
Rings

- Appear in most CT acquisitions
- Caused by stuck pixels in the projection data
- Can mostly be suppressed during reconstruction



Lines

- Frequent in neutron CT slices
- Caused by spots on single projections
- Can mostly be suppressed during reconstruction

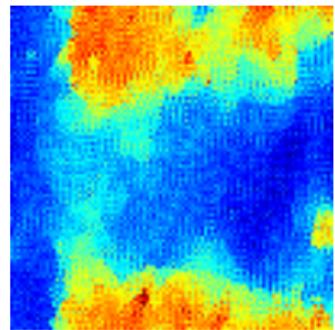


Rounding errors

- May appear with sum operations on large data sets.
- At some point the new term is smaller than the precision.

Instable processing

- Due to incorrect regularization
- Wrong parameterization
- Incorrect implementation... bugs etc



rand

Generates an $m \times n$ random field with uniform distribution

randn

Generates an $m \times n$ random field with Gaussian distribution

poissrnd

Generates an $m \times n$ random field with Poisson distribution

mean, var, std

Computes the mean, variance, and standard deviation of an image f

min,max

Computes the min and max value of an image f

Basic filtering

The first approach to image enhancement

General definition

A filter is a processing unit that

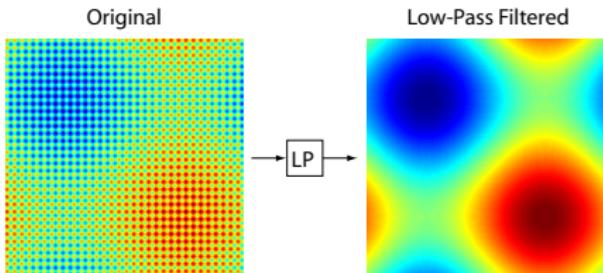
- Enhances the wanted information
- Suppresses the unwanted information

Ideally without altering relevant features beyond recognition

[Jähne, 2005]

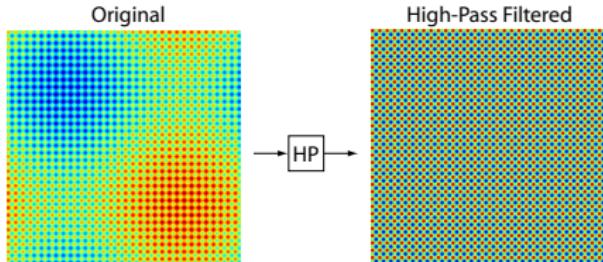
Low pass filters

- Slow changes are enhanced
- Rapid changes are suppressed



High pass filters

- Rapid changes are enhanced
- Slow changes are suppressed



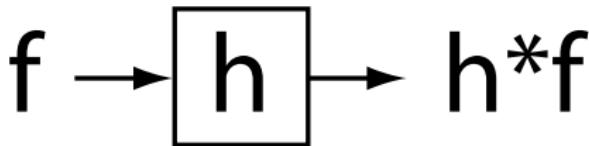
Linear filters

Computed using the convolution operation

$$g(x) = h * f(x) = \int_{\Omega} f(x - \tau) h(\tau) d\tau \quad (3)$$

where

- $f(x)$ is the image
- h is the convolution kernel of the filter



Low-pass filter kernels

Mean or Box filter

All weights have the same value.

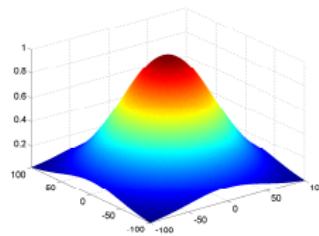
Example:

$$B = \frac{1}{25} \cdot \begin{array}{|c|c|c|c|c|} \hline 1 & 1 & 1 & 1 & 1 \\ \hline 1 & 1 & 1 & 1 & 1 \\ \hline 1 & 1 & 1 & 1 & 1 \\ \hline 1 & 1 & 1 & 1 & 1 \\ \hline 1 & 1 & 1 & 1 & 1 \\ \hline \end{array}$$

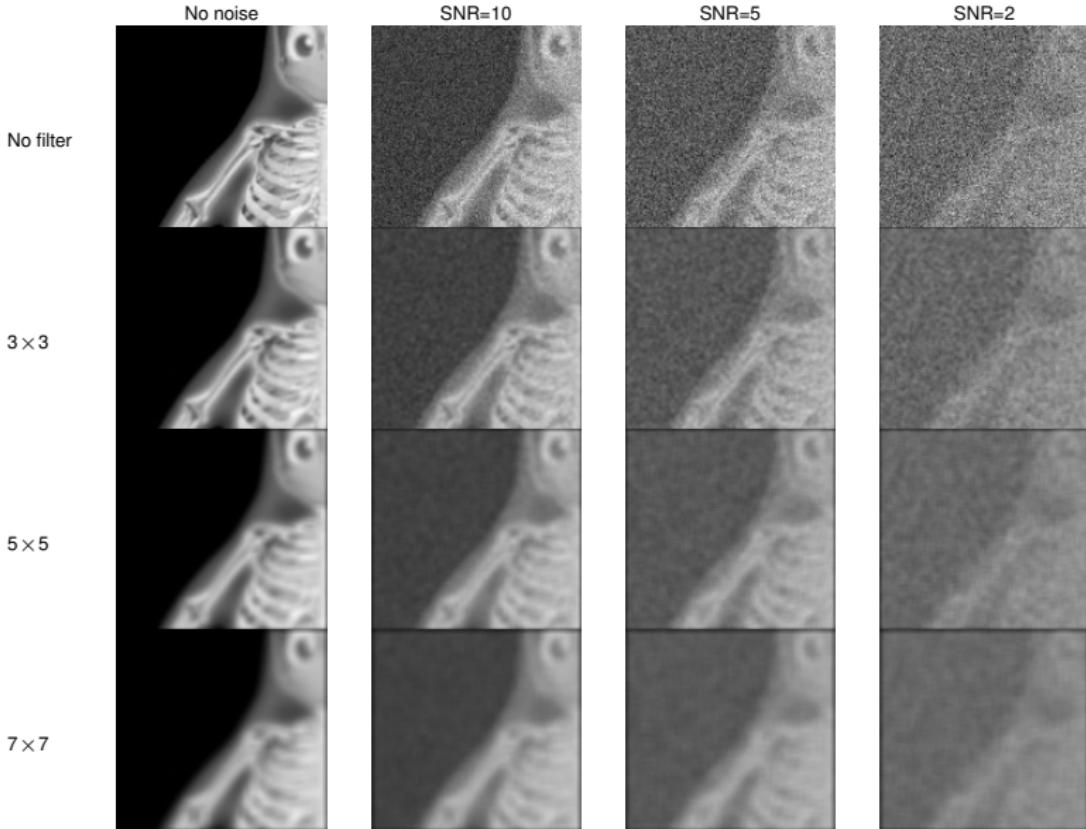
Gauss filter

$$G = e^{-\frac{x^2+y^2}{2\sigma^2}}$$

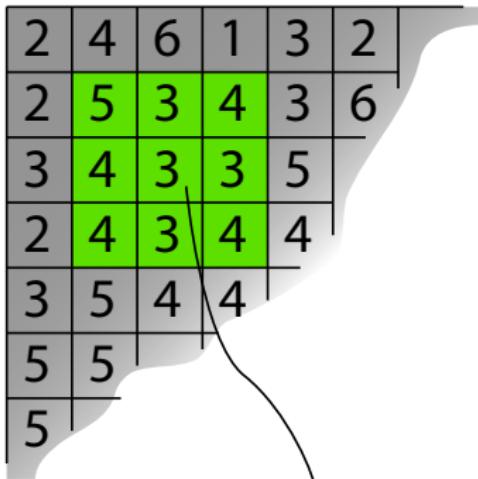
Example:



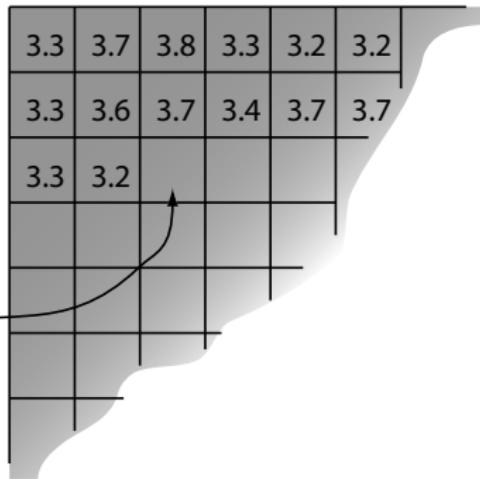
Using a Mean filter



How is the convolution computed



$$(5+3+4+4+3+3+4+3+4)/9=3.7$$



Note

For a non-uniform kernel each term is weighted by its kernel weight.

The associative and commutative laws apply to convolution

$$(a * b) * c = a * (b * c) \quad \text{and} \quad a * b = b * a$$

A convolution kernel is called *separable* if it can be split in two or more parts:

$$\begin{array}{|c|c|c|} \hline \cdot & \cdot & \cdot \\ \hline \cdot & \cdot & \cdot \\ \hline \cdot & \cdot & \cdot \\ \hline \end{array} = \begin{array}{|c|} \hline \cdot \\ \hline \cdot \\ \hline \cdot \\ \hline \end{array} * \begin{array}{|c|c|c|} \hline \cdot & \cdot & \cdot \\ \hline \end{array}$$

Gain

Reduces the number of computations \rightarrow faster processing

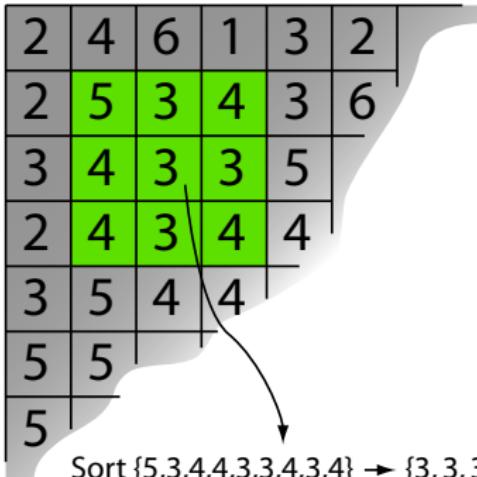
$$3 \times 3 \rightarrow 9 \text{ mult and } 8 \text{ add} \Leftrightarrow 6 \text{ mult and } 4 \text{ add}$$

$$3 \times 3 \times 3 \rightarrow 27 \text{ mult and } 26 \text{ add} \Leftrightarrow 9 \text{ mult and } 6 \text{ add}$$

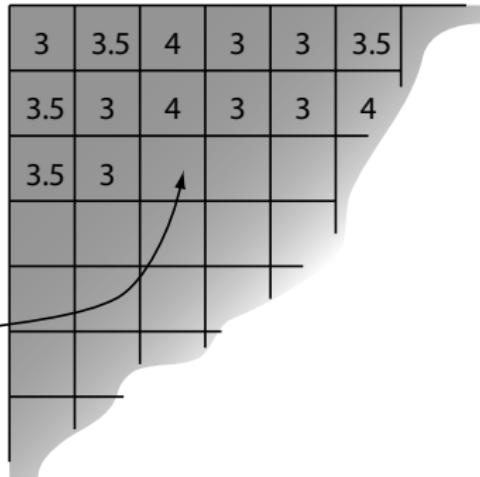
Example

$$e^{-\frac{x^2+y^2}{2\sigma^2}} = e^{-\frac{x^2}{2\sigma^2}} * e^{-\frac{y^2}{2\sigma^2}}$$

The median filter



Sort $\{5, 3, 4, 4, 3, 3, 4, 3, 4\} \rightarrow \{3, 3, 3, 3, 4, 4, 4, 4, 5\}$



Comparing filters for different noise types

10% salt&pepper noise



Median filtered



Mean filtered



Additive White Gaussian noise
($\sigma = 30$)



Median filtered



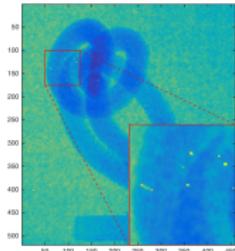
Mean filtered



Filter example: Spot cleaning

Problem

- Many neutron images are corrupted by spots that confuse following processing steps.
- The amount, size, and intensity varies with many factors.

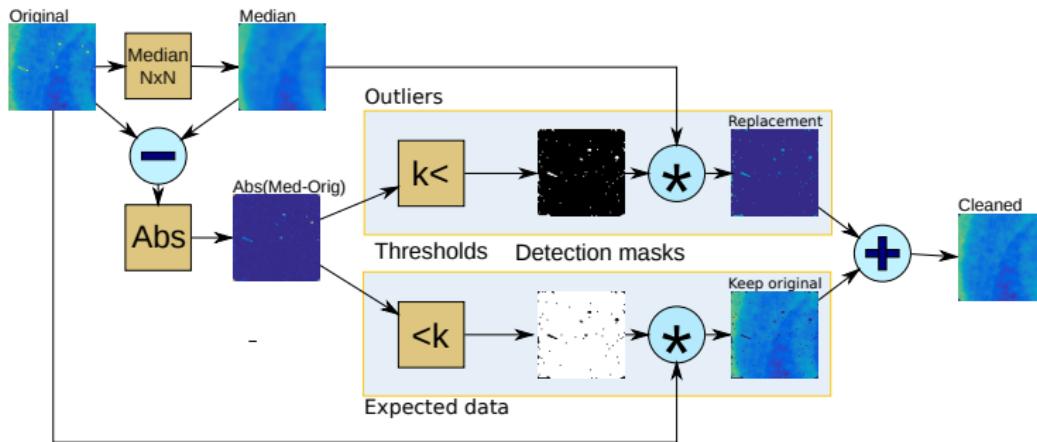


Solutions

- :-) Low pass filter
- :-) Median filter
- :-) Detect spots and replace by estimate

Spot cleaning algorithm

An algorithm

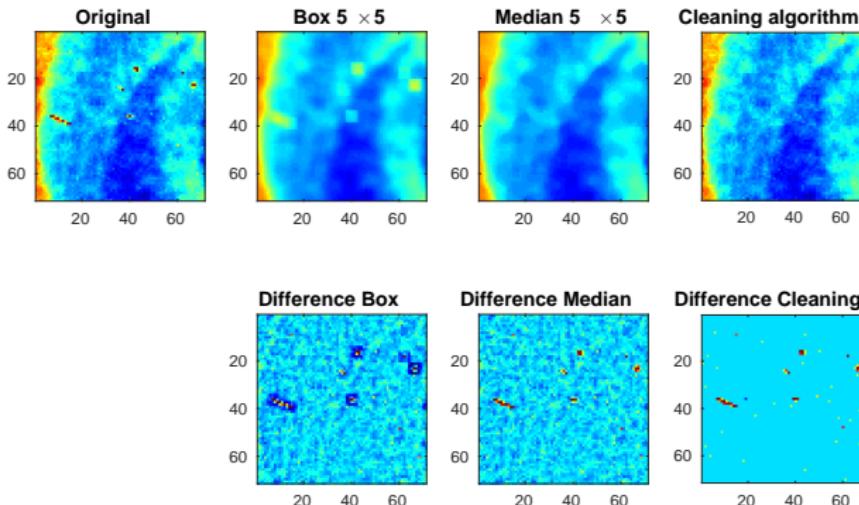


Parameters

N Width of median filter.

k Threshold level for outlier detection.

Spot cleaning – Compare performance



The ImageJ ways

Despeckle Median ... please avoid this one!!!

Remove outliers Similar to cleaning described algorithm

High-pass filters

High-pass filters enhance rapid changes – ideal for edge detection

Typical high-pass filters:

Gradients

$$\frac{\partial}{\partial x} = \frac{1}{2} \cdot \begin{bmatrix} -1 & 1 \end{bmatrix}$$

Laplacian

$$\Delta = \frac{1}{2} \cdot \begin{bmatrix} 1 & 2 & 1 \\ 2 & -12 & 2 \\ 1 & 2 & 1 \end{bmatrix}$$

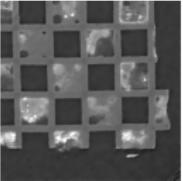
$$\frac{\partial}{\partial x} = \frac{1}{32} \cdot \begin{bmatrix} -3 & 0 & 3 \\ -10 & 0 & 10 \\ -3 & 0 & 3 \end{bmatrix}$$

Sobel

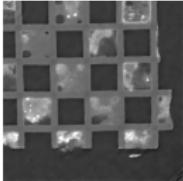
$$G = |\nabla f| = \sqrt{\left(\frac{\partial}{\partial x}f\right)^2 + \left(\frac{\partial}{\partial y}f\right)^2}$$

Gradient example

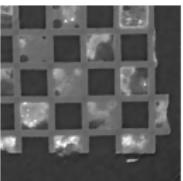
Vertical edges

$$\frac{\partial}{\partial x} \begin{matrix} \text{Image} \\ \text{(Input)} \end{matrix} = \frac{1}{32} \cdot \begin{matrix} \text{Kernel} \\ \begin{array}{|c|c|c|} \hline -3 & 0 & 3 \\ \hline -10 & 0 & 10 \\ \hline -3 & 0 & 3 \\ \hline \end{array} \end{matrix} * \begin{matrix} \text{Image} \\ \text{(Input)} \end{matrix} = \begin{matrix} \text{Image} \\ \text{(Output)} \end{matrix}$$


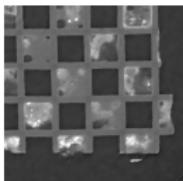
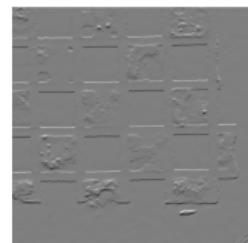
-3	0	3
-10	0	10
-3	0	3




Horizontal edges

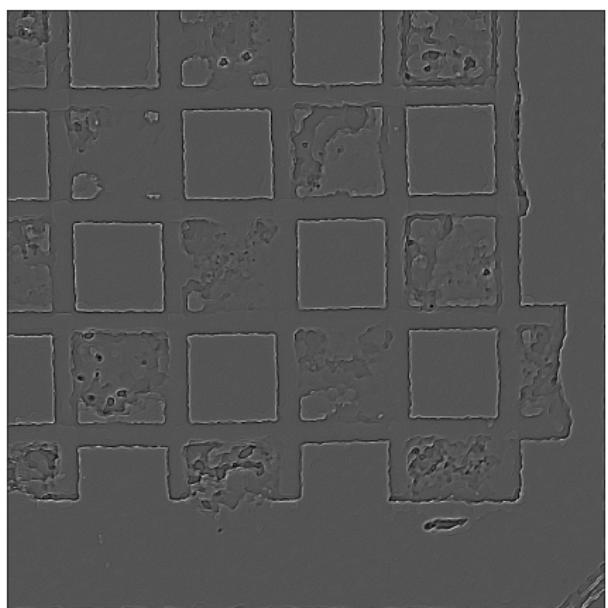
$$\frac{\partial}{\partial y} \begin{matrix} \text{Image} \\ \text{(Input)} \end{matrix} = \frac{1}{32} \cdot \begin{matrix} \text{Kernel} \\ \begin{array}{|c|c|c|} \hline -3 & -10 & -3 \\ \hline 0 & 0 & 0 \\ \hline 3 & 10 & 3 \\ \hline \end{array} \end{matrix} * \begin{matrix} \text{Image} \\ \text{(Input)} \end{matrix} = \begin{matrix} \text{Image} \\ \text{(Output)} \end{matrix}$$


-3	-10	-3
0	0	0
3	10	3

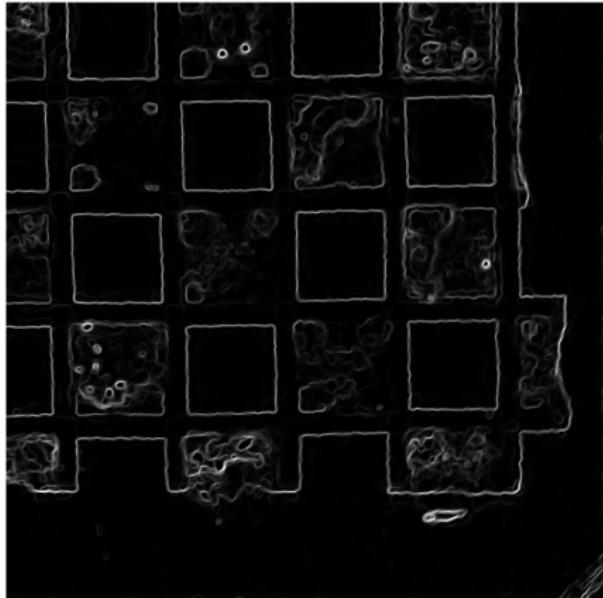
Edge detection examples

Laplacian



Both negative and positive values

Sobel



Positive values only.

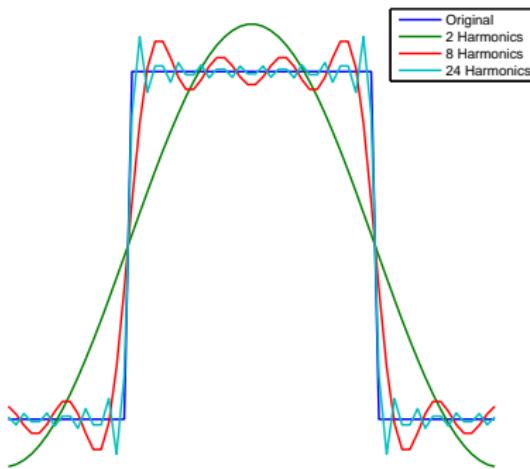
Filters in frequency space

Applications of the Fourier transform

What is frequency space?

Introduction

- A signal can be decomposed into a sum of basic harmonics defined by amplitude, phase shift and frequency.
- Fine details and sharp edges require more harmonics



Transform

$$G(\xi_1, \xi_2) = \mathcal{F}\{g\} = \int_{-\infty}^{\infty} \int_{-\infty}^{\infty} g(x, y) e^{-i(\xi_1 x + \xi_2 y)} dx dy$$

It's inverse

$$g(x, y) = \mathcal{F}^{-1}\{G\} = \frac{1}{(2\pi)^2} \int_{-\infty}^{\infty} \int_{-\infty}^{\infty} G(\omega) e^{i(\xi_1 x + \xi_2 y)} d\xi_1 d\xi_2$$

FFT

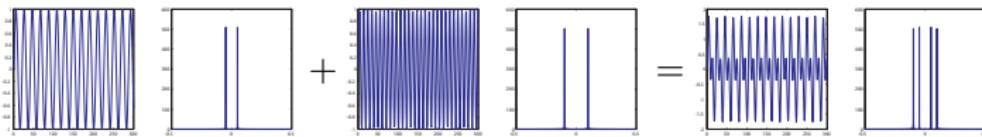
In practice – you never see the transform equations.

The Fast Fourier Transform is available in numerical libraries and tools.

[Jähne, 2005]

Addition

$$\mathcal{F}\{a + b\} = \mathcal{F}\{a\} + \mathcal{F}\{b\}$$

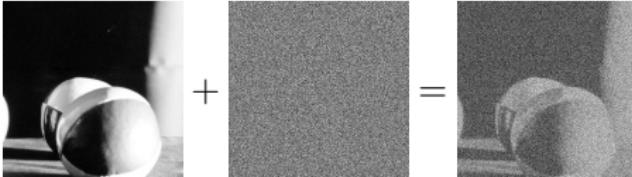


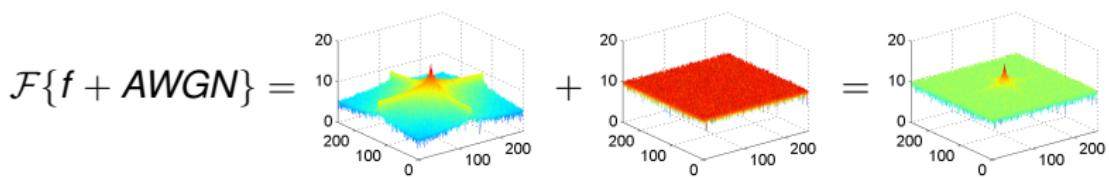
Convolution

$$\mathcal{F}\{a * b\} = \mathcal{F}\{a\} \cdot \mathcal{F}\{b\}$$

$$\mathcal{F}\{a \cdot b\} = \mathcal{F}\{a\} * \mathcal{F}\{b\}$$

Additive noise in Fourier space

$$f + AWGN = \text{Image} + \text{Noise} = \text{Result}$$




Problem

How can we suppress noise without destroying relevant image features?

0°

$$\mathcal{F} \left\{ \begin{array}{c} \text{[Image of a vertical grid]} \\ \end{array} \right\} \Rightarrow \text{[Image of a blue rectangle]} + \text{[Image of a small white dot]}$$

90°

$$\mathcal{F} \left\{ \begin{array}{c} \text{[Image of a stack of horizontal bars]} \end{array} \right\} \Rightarrow \begin{array}{c} \text{[Image of a single blue dot]} \end{array}$$

30°

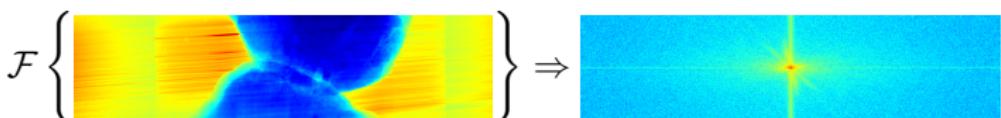
$$\mathcal{F} \left\{ \begin{array}{c} \text{hatched image} \end{array} \right\} \Rightarrow \begin{array}{c} \text{blue image} \end{array}$$

60°

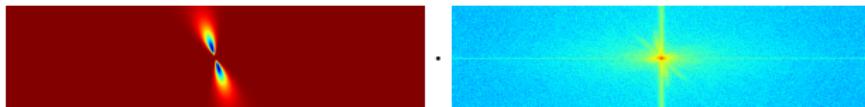
$$\mathcal{F} \left\{ \begin{array}{c} \text{[diagonal hatching]} \\ \end{array} \right\} \Rightarrow \begin{array}{c} \text{[blue square]} \\ \text{[dots]} \end{array}$$

Example – Stripe removal in Fourier space

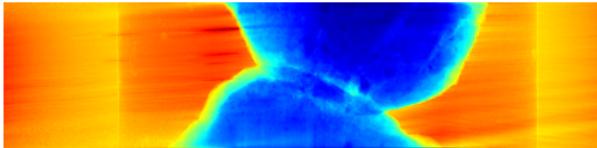
- 1 Transform the image to Fourier space



- 2 Multiply spectrum image by band pass filter

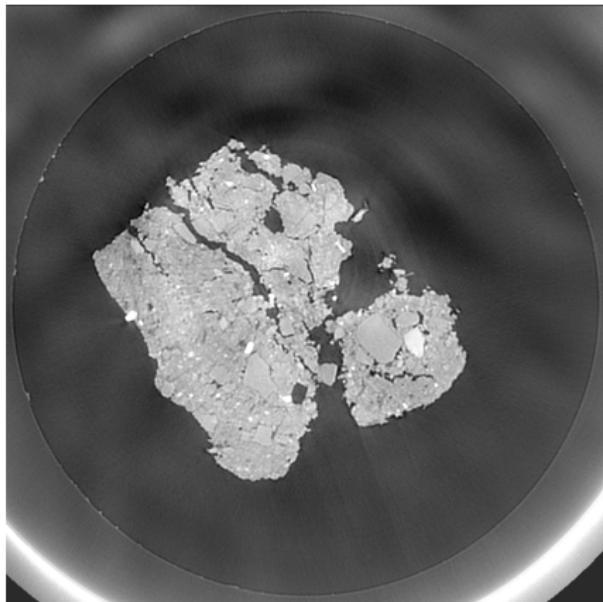


- 3 Compute the inverse transform to obtain the filtered image in real space

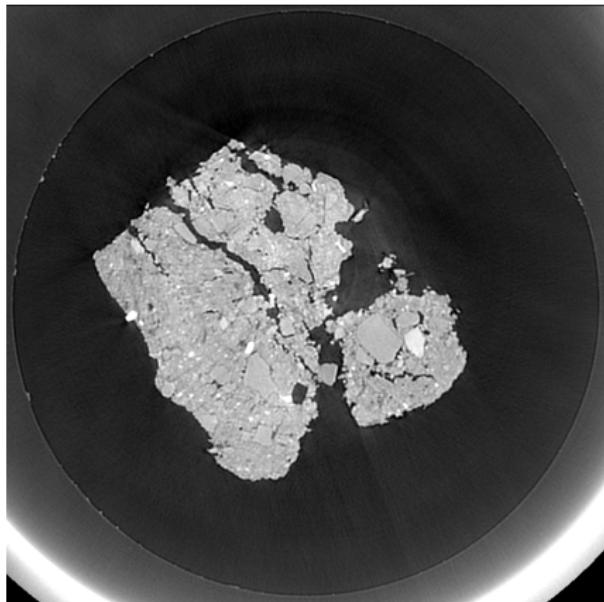


The effect of the stripe filter

Reconstructed CT slice before filter



Reconstructed CT slice after filter



Intensity variations are suppressed using the stripe filter on all projections.

filter2(h,f)

Linear filter using kernel h on image f .

medfilt2(f,[x y])

Median filter using an $x \times y$ filter neighborhood

fft2(f), ifft2(F)

Computes the Fast Fourier Transform and its inverse of image f .

abs(f), angle(f)

Computes amplitude and argument of a complex number.

real(f), imag(f)

Gives the real and imaginary parts of a complex number.

Scale spaces

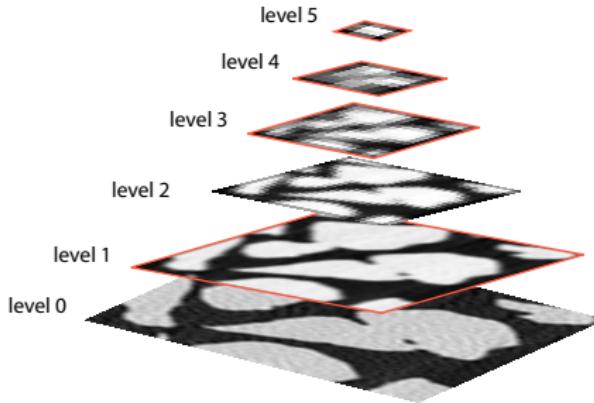
Why scale spaces?

Motivation

Basic filters have problems to handle low SNR and textured noise.
Something new is required...

The solution

Filtering on different scales can take noise suppression one step further.



Wavelets – the basic idea

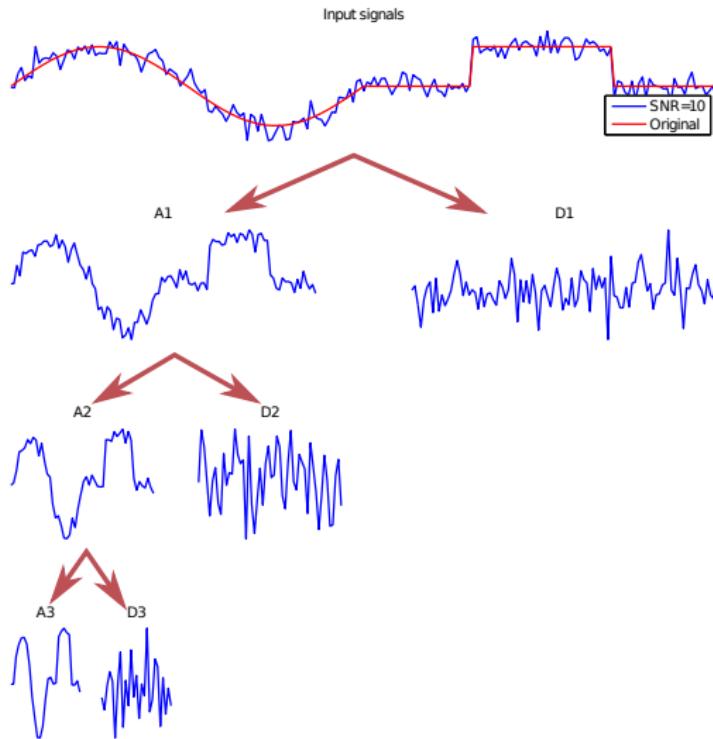
- The wavelet transform produces scales by decomposing a signal into two signals at a coarser scale containing *trend* and *details*.
- The next scale is computed using the trend of the previous transform

$$WT\{s\} \rightarrow \{a_1, d_1\}, WT\{a_1\} \rightarrow \{a_2, d_2\}, \dots, WT\{a_{N-1}\} \rightarrow \{a_N, d_N\}$$

- The inverse transform brings s back using $\{a_N, d_1, \dots, d_N\}$.
- Many wavelet bases exists, the choice depends on the application.
- Wavelets can have several uses:
 - Noise reduction
 - Analysis
 - Segmentation
 - Compression

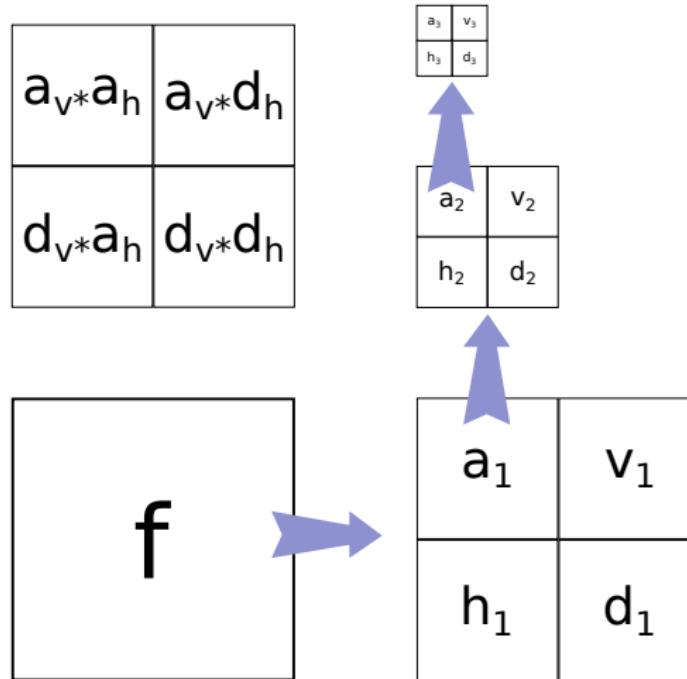
[J.C.Walker, 1999][Mallat, 2009]

Wavelet transform of a 1D signal



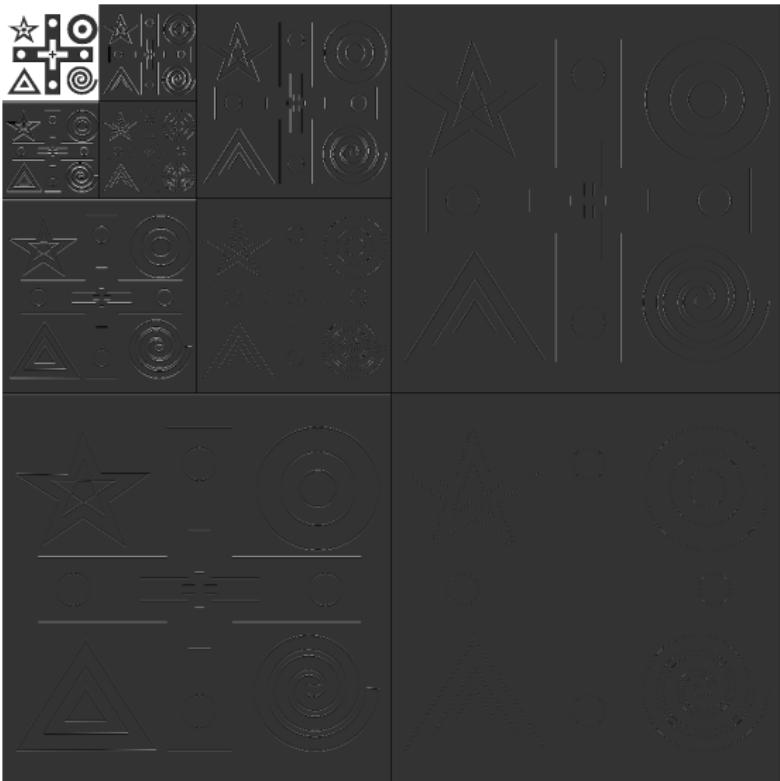
Using *symlet-4*

Wavelet transform of an image



Wavelet transform of an image – example

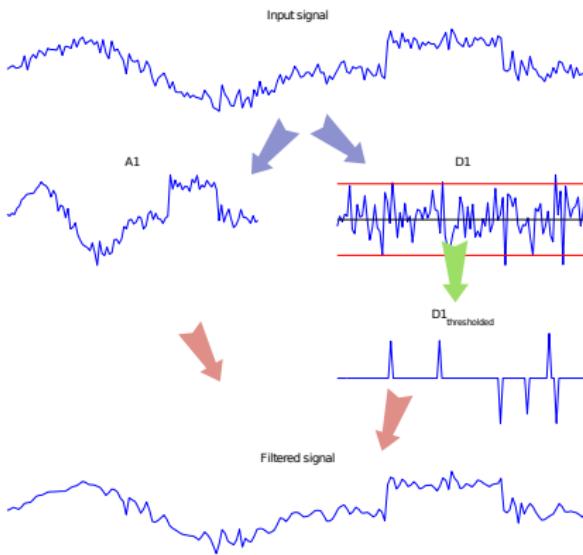
$$WT \left\{ \begin{array}{c} \star \\ \square \\ \triangle \end{array} \right. \left\{ \begin{array}{c} \bullet \\ + \\ \circlearrowleft \end{array} \right. =$$



Wavelet noise reduction

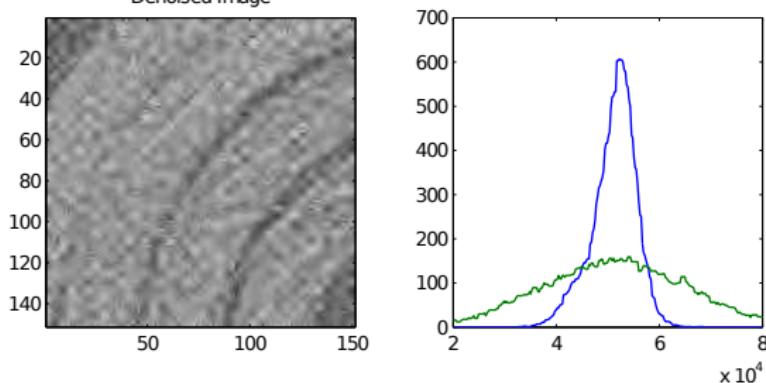
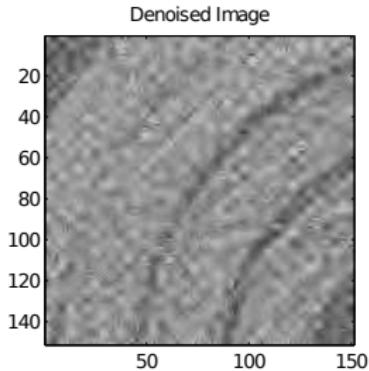
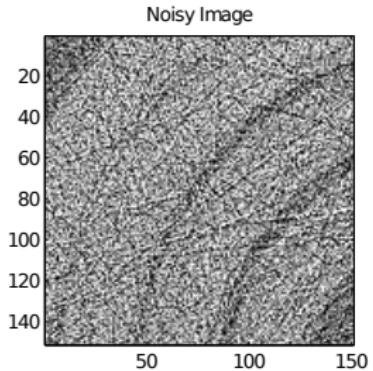
The noise is found in the detail part of the WT

- Make a WT of the signal to a level that corresponds to the scale of the unwanted information.
- Threshold the detail part $d_\gamma = |d| < \gamma ? 0 : d$.
- Inverse WT back to normal scale → image is filtered.



Wavelet noise reduction – Image example

Filtered using two levels of Symlet-2 wavelet



Neutron CT of a lead scroll.

dwt2/idtw2

Makes one level of the wavelet transform or its inverse using wavelet base specified by 'wn'.

wavedec2

Performs N levels of wavelet decomposition using a specified wavelet base.

wbmpen

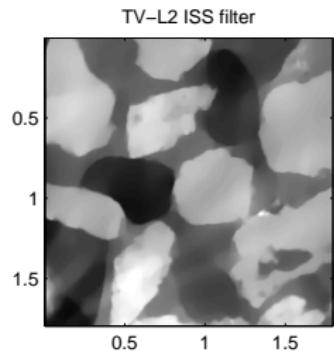
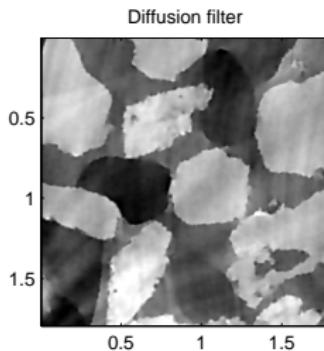
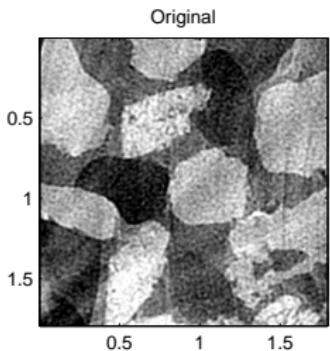
Estimating threshold parameters for wavelet denoising.

wdencmp

Wavelet denoising and compression using information from *wavedec2* and *wbmpen*.

PDE based scale space filters

Filters small features faster than larger ones.



May work for applications where Linear and Rank filters fail.

[Aubert and Kornprobst, 2002].

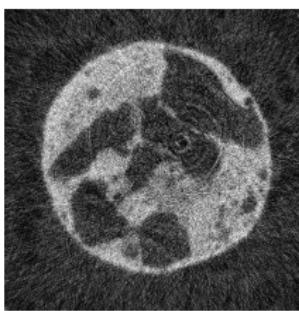
The starting point

The heat transport equation

$$\frac{\partial T}{\partial t} = \kappa \nabla^2 T$$

T Image to filter (intensity \equiv temperature)

κ Thermal conduction capacity



Original

Intensity diffusion

The steady state solution is a homogeneous image...

Controlling the diffusivity

We want to control the diffusion process...

Near edges The Diffusivity $\rightarrow 0$

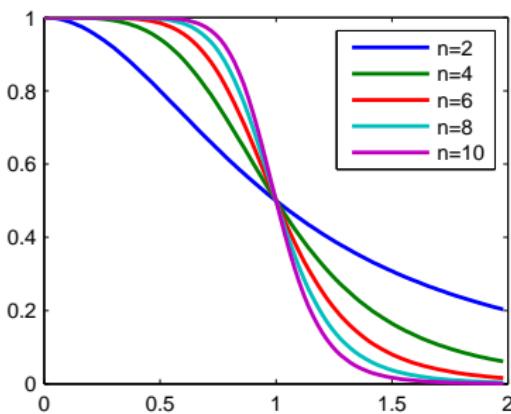
Flat regions The Diffusivity $\rightarrow 1$

The contrast function G is our control function

$$G(x) = \frac{1}{1 + \left(\frac{x}{\lambda}\right)^n}$$

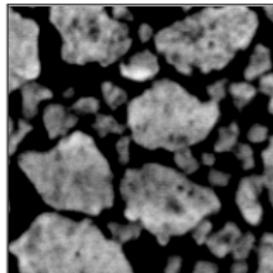
λ Threshold level

n Steepness of the threshold

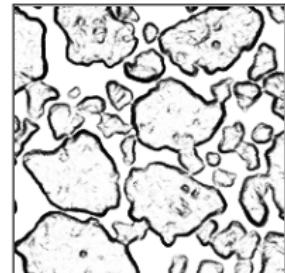


Gradient controlled diffusivity

$$\frac{\partial u}{\partial t} = G(|\nabla u|) \nabla^2 u$$



Image



Diffusivity map

u Image to be filtered

G(·) Non-linear function to control the diffusivity

τ Time increment

N Number of iterations

This filter is noise sensitive!

A more robust filter is obtained with

$$\frac{\partial u}{\partial t} = G(|\nabla_\sigma u|) \nabla^2 u \quad (4)$$

u Image to be filtered

G(·) Non-linear function to control the contrast

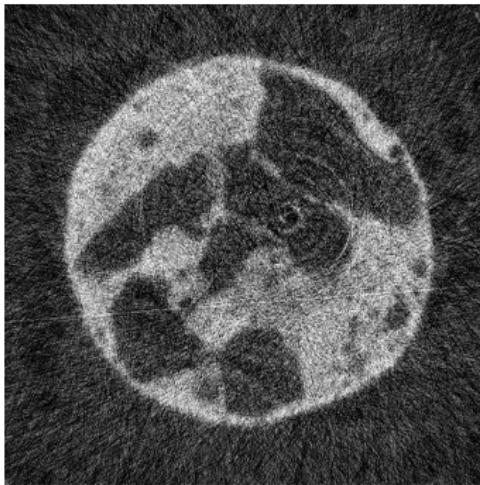
τ Time increment per numerical iteration

N Number of iterations

∇_σ Gradient smoothed by a Gaussian filter, width σ

Diffusion filter example

Neutron CT slice from a real-time experiment observing the coalescence of cold mixed bitumen.



Original

Iterations of non-linear diffusion

The continued development

- 90's** During the late 90's the diffusion filter was described in terms of a regularization problem.
- 00's** Work toward regularization of total variation minimization.

TV-L¹

$$u = \operatorname{argmin}_{u \in BV(\Omega)} \left\{ \underbrace{|u|_{BV}}_{\text{noise}} + \underbrace{\frac{\lambda}{2} \|f - u\|_1}_{\text{fidelity}} \right\}$$

Rudin-Osher-Fatemi model (ROF)

$$u = \operatorname{argmin}_{u \in BV(\Omega)} \left\{ \underbrace{|u|_{BV}}_{\text{noise}} + \underbrace{\frac{\lambda}{2} \|f - u\|_2^2}_{\text{fidelity}} \right\}$$

with $|u|_{BV} = \int_{\Omega} |\nabla u|^2$

The idea

We want smooth regions with sharp edges...

- Turn the processing order of scale space filter upside down
- Start with an empty image
- Add large structures successively until an image with relevant features appears

The ISS filter – Some properties

- is an edge preserving filter for noise reduction.
- is defined by a partial differential equation.
- has a well defined termination point.

[Burger et al., 2006]

The ROF filter equation

The image f is filtered by solving

$$\begin{aligned}\frac{\partial u}{\partial t} &= \operatorname{div} \left(\frac{\nabla u}{|\nabla u|} \right) + \lambda (f - u + v) \\ \frac{\partial v}{\partial t} &= \alpha (f - u)\end{aligned}\tag{5}$$

Variables:

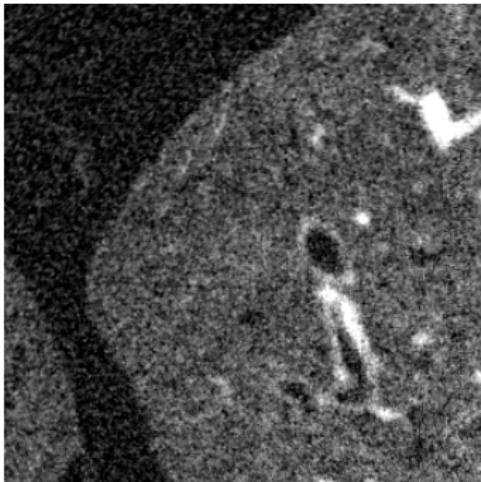
- f Input image
- u Filtered image
- v Regularization term (feedback of previous iteration)

Filter parameters

- λ Related to the scale of the features to suppress.
- α Quality refinement
- N Number of iterations
- τ Time increment

Filter iterations

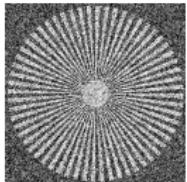
Neutron CT of dried lung filtered using 3D ISS filter



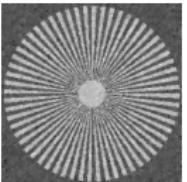
Original

Filter iterations

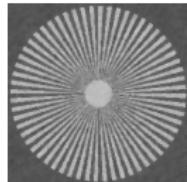
Solutions at different times



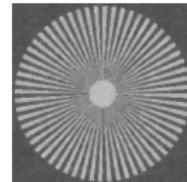
1



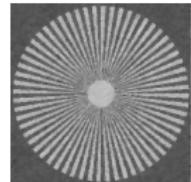
30



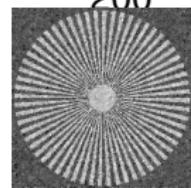
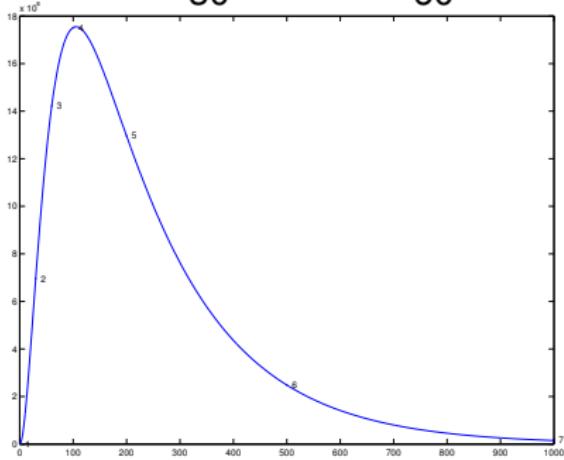
60



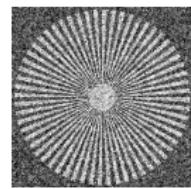
100



200



500



999

The idea

Smoothing normally consider information from the neighborhood like

- Local averages (convolution)
- Gradients and Curvatures (PDE filters)

Non-local smoothing average similiar intensities in a global sense.

- Every filtered pixel is a weighted average of all pixels.
- Weights computed using difference between pixel intensities.

[Buades et al., 2005]

Filter definition

The non-local means filter is defined as

$$u(p) = \frac{1}{C(p)} \sum_{q \in \Omega} v(q) f(p, q) \quad (6)$$

where

v and u input and result images.

$C(p)$ is the sum of all pixel weights as

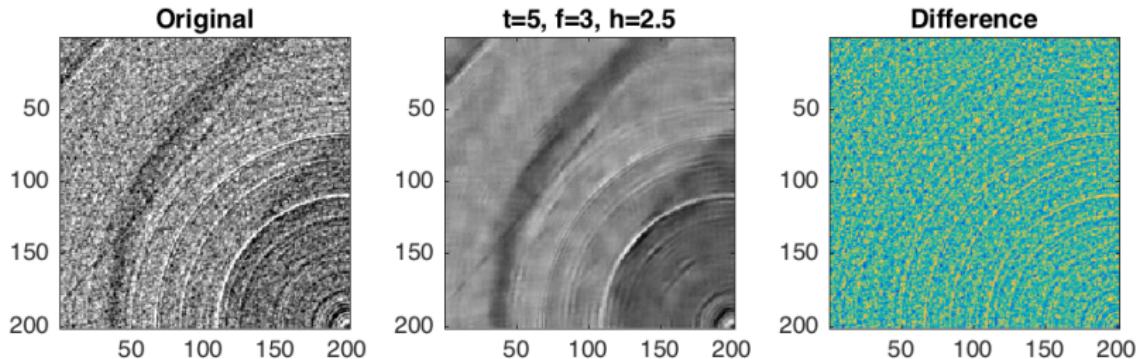
$$C(p) = \sum_{q \in \Omega} f(p, q) \quad (7)$$

$f(p, q)$ is the weighting function

$$f(p, q) = e^{-\frac{|B(q) - B(p)|^2}{h^2}} \quad (8)$$

$B(x)$ is a neighborhood operator e.g. local average around x

Non-local means 2D – Example



Observations

- Good smoothing effect.
- Strong thin lines are preserved.
- Some patchiness related to filter parameter t , i.e. the size of Ω_i .

Problem

The original filter compares all pixels with all pixels...

- Complexity $\mathcal{O}(N^2)$
- Not feasible for large images, and particular 3D images!

Solution

It has been shown that not all pixels have to be compared to achieve a good filter effect.

i.e. Ω in eq 6 and 7 can be replaced by $\Omega_i \ll \Omega$

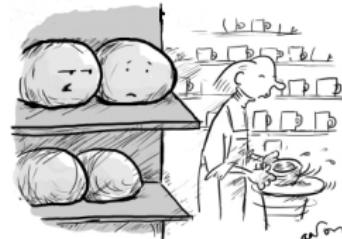
Verification

Verify the correctness of the method

"Data massage"

Filtering manipulates the data...

... avoid too strong modifications otherwise
you may invent new image features!!!



Watch that man, he'll make mugs of us all!

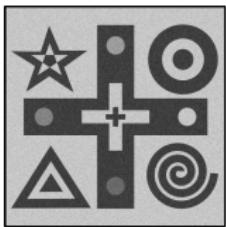
Verify the validity your method

- Visual inspection
- Difference images
- Use degraded phantom images in a "smoke test"

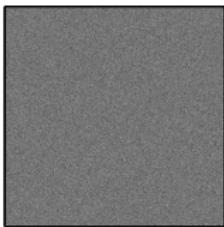
Verification using difference images

Compute pixel-wise difference between image f and g

Noisy image



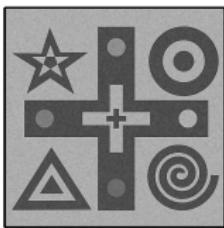
Ideal filter



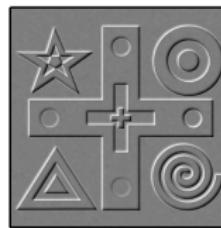
Over smoothing



Intensity scaling



Geometric shift



Difference images provide first diagnosis about processing performance

- Testing term from electronic hardware testing – drive the system until something fails due to overheating...
- In general: scan the parameter space for different SNR until the method fails to identify strength and weakness of the system.

Test strategy

- 1 Create a phantom image with relevant features.
- 2 Add noise for different SNR to the phantom.
- 3 Apply the processing method with different parameters.
- 4 Measure the difference between processed and phantom.
- 5 Repeat steps 2-4 N times for better test statistics.
- 6 Plot the results and identify the range of SNR and parameters that produce acceptable results.

An evaluation procedure need a metric to compare the performance

Mean squared error

$$MSE(f, g) = \sum_{p \in \Omega} (f(p) - g(p))^2$$

Structural similarity index

$$SSIM(f, g) = \frac{(2\mu_f \mu_g + C_1)(2\sigma_{fg} + C_2)}{(\mu_f^2 + \mu_g^2 + C_1)(\sigma_f^2 + \sigma_g^2 + C_2)}$$

μ_f, μ_g Local mean of f and g .

σ_{fg} Local correlation between f and g .

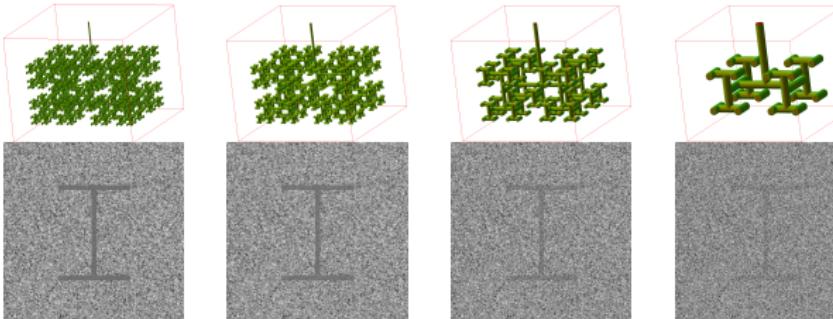
σ_f, σ_g Local standard deviation of f and g .

C_1, C_2 Constants based on the image dynamics (small numbers).

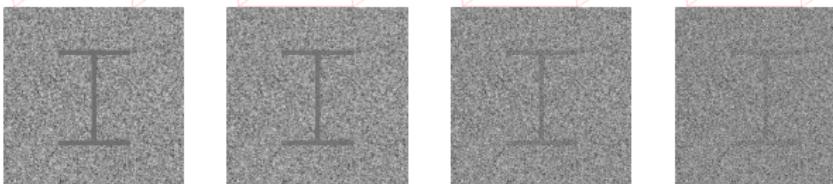
$$MSSIM(f, g) = E[SSIM(f, g)]$$

Test run example

Phantom structures

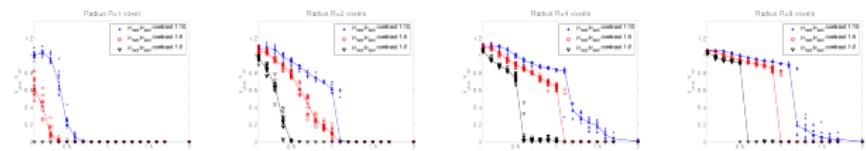


Add noise



Process

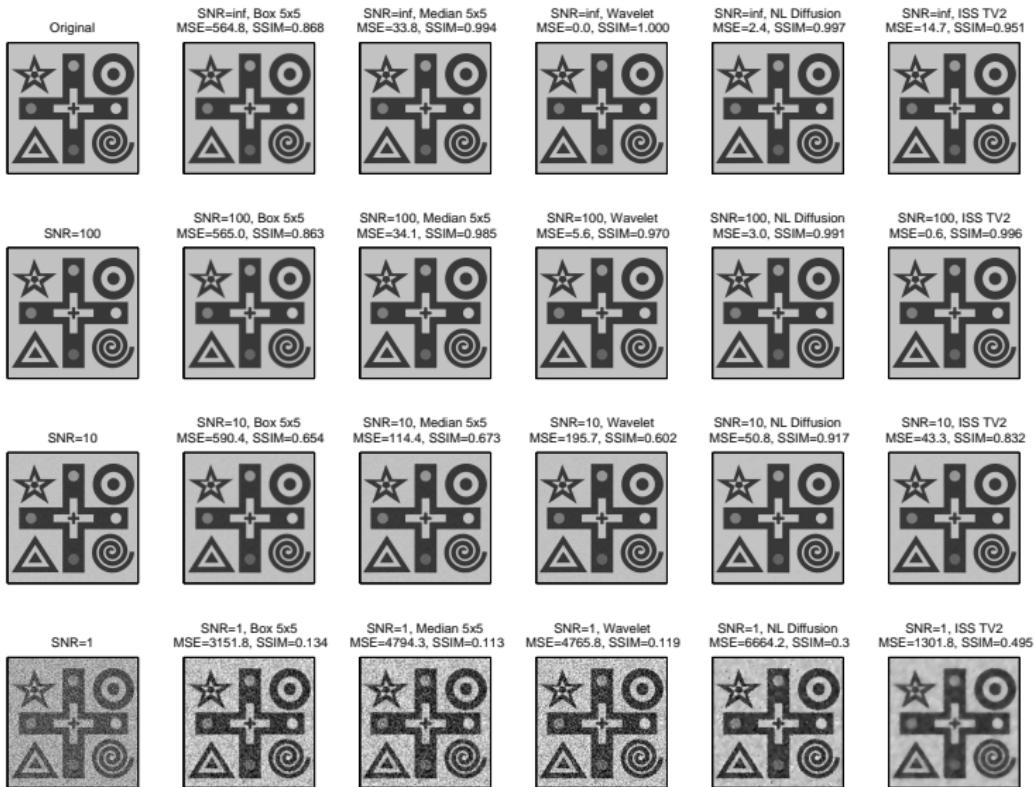
Plot results



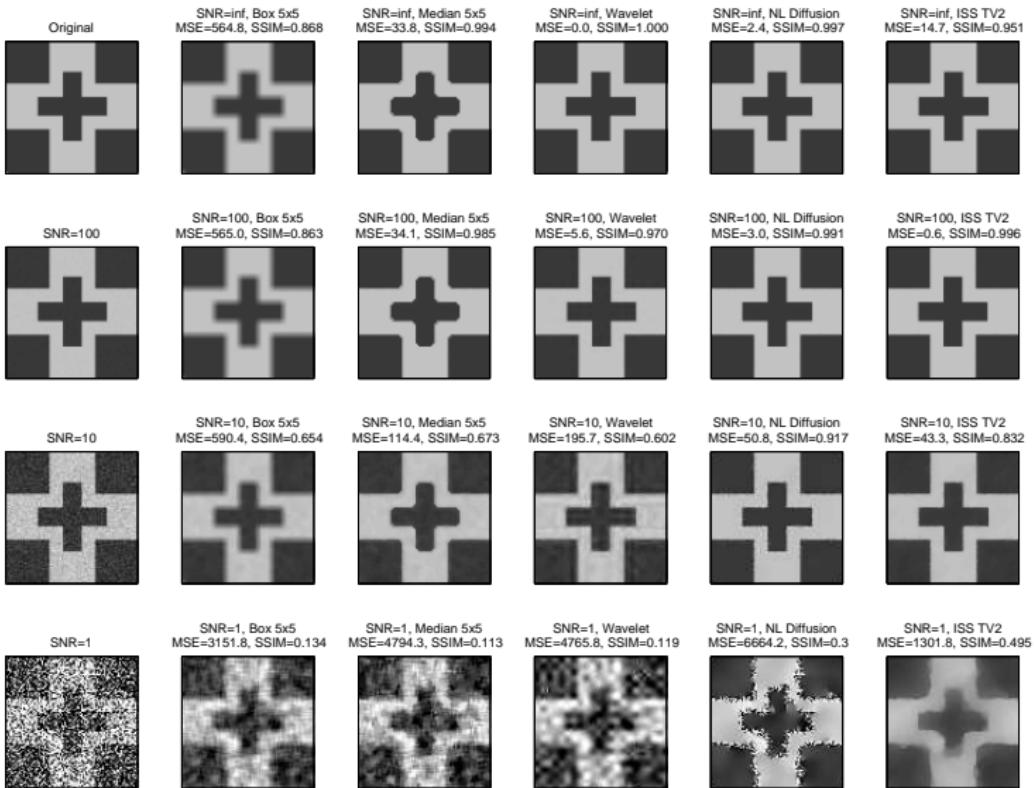
[Kaestner et al., 2006]

Summary

Many filters



Details of filter performance



Take-home message

We have looked at different ways to suppress noise and artifacts:

- Convolution
- Median filters
- Wavelet denoising
- PDE filters

Which one you select depends on

- Purpose of the data
- Quality requirements
- Available time

Remember

A good measurement is better than an enhanced bad measurement...
...but bad data can mostly be rescued if needed.

References I



Aubert, G. and Kornprobst, P. (2002).

Mathematical problems in image processing.

Number 147 in Applied mathematical sciences. Springer Verlag.



Buades, A., Coll, B., and Morel, J. (2005).

A non-local algorithm for image denoising.

In *IEEE Int. Conf. on Computer Vision and Pattern Recognition*.



Burger, M., Gilboa, G., Osher, S., and Xu, J. (2006).

Nonlinear inverse scale space methods.

Communications in Mathematical Sciences, 4(1):179–212.



Jähne, B. (2005).

Digital Image Processing.

Springer Verlag, 6 edition.



J.C. Walker (1999).

A primer on Wavelets and their scientific applications.

CRC Press.



Kaestner, A., Schneebeli, M., and Graf, F. (2006).

Visualizing three-dimensional root networks using computed tomography.

Geoderma, 136(1–2):459–469.



Mallat, S. (2009).

A wavelet tour of signal processing: The sparse way.

Academic Press.



Wang, Z. and Bovik, A. (2009).

Mean squared error: Love it or leave it? – a new look at signal fidelity measures.

IEEE Signal Processing Magazine, January:89–117.